

customer_360_clean

August 21, 2025

1 Customer 360 - Project Context

This Customer-360 extends our earlier work—same cleaned data foundation, now focused on retention & personalization to fix the leaky bucket.

1. Findings from previous analysis:

- Stagnation (−1.2% YoY in 2011)
- 37% new vs 37% lost customers
- Revenue dependence on loyalists (~44% of base → ~88% of revenue)

2 Project Setup

All libraries imported successfully!

Dataset shape: (732583, 15)

Date range: 2009-12-01 07:45:00 to 2011-12-09 12:50:00

Number of unique customers: 5819

Number of unique products: 4604

Total revenue: \$14,575,657.25

	order_id	product_id	product_description	quantity	\
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	
1	489434	79323P	PINK CHERRY LIGHTS	12	
2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	

	order_date	unit_price	customer_id	country	total_amount	\
0	2009-12-01 07:45:00	6.95	13085.0	United Kingdom	83.4	
1	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.0	
2	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.0	
3	2009-12-01 07:45:00	2.10	13085.0	United Kingdom	100.8	
4	2009-12-01 07:45:00	1.25	13085.0	United Kingdom	30.0	

	year	month	quarter	day_of_week	month_year	product_category
0	2009	12	4	Tuesday	2009-12	CHRISTMAS_HOLIDAY
1	2009	12	4	Tuesday	2009-12	BEAUTY_PERSONAL
2	2009	12	4	Tuesday	2009-12	BEAUTY_PERSONAL

3	2009	12	4	Tuesday	2009-12	HOME_DECOR
4	2009	12	4	Tuesday	2009-12	HOME_DECOR

We have cleaned the dataset in the separate analysis, so we expect this dataset doesn't need extensive preprocessing anymore

=== DATA QUALITY ASSESSMENT ===

1. Missing Values:

order_id	0
product_id	0
product_description	0
quantity	0
order_date	0
unit_price	0
customer_id	0
country	0
total_amount	0
year	0
month	0
quarter	0
day_of_week	0
month_year	0
product_category	0
dtype:	int64

2. Data Types:

order_id	int64
product_id	object
product_description	object
quantity	int64
order_date	object
unit_price	float64
customer_id	float64
country	object
total_amount	float64
year	int64
month	int64
quarter	int64
day_of_week	object
month_year	object
product_category	object
dtype:	object

3. Duplicate Records:

Total duplicates: 0

4. Customer ID Analysis:

Missing customer IDs: 0
Customer ID data type: float64

5. Data Anomalies:
Negative quantities: 0
Negative unit prices: 0
Negative total amounts: 0

6. Basic Statistics:

	order_id	quantity	unit_price	customer_id \
count	732583.000000	732583.000000	732583.000000	732583.000000
mean	537582.753560	12.523575	2.865135	15336.361872
std	26948.060684	71.180639	3.934765	1691.785784
min	489434.000000	1.000000	0.030000	12346.000000
25%	514531.000000	2.000000	1.250000	13985.000000
50%	536856.000000	5.000000	1.950000	15281.000000
75%	562269.000000	12.000000	3.750000	16805.500000
max	581587.000000	19152.000000	295.000000	18287.000000

	total_amount	year	month	quarter
count	732583.000000	732583.000000	732583.000000	732583.000000
mean	19.896254	2010.433758	7.438405	2.804439
std	60.314327	0.568138	3.419965	1.131955
min	0.060000	2009.000000	1.000000	1.000000
25%	4.950000	2010.000000	5.000000	2.000000
50%	11.550000	2010.000000	8.000000	3.000000
75%	18.240000	2011.000000	11.000000	4.000000
max	8925.000000	2011.000000	12.000000	4.000000

3 Phase 1: Creating a single customer view aggregation

In this phase, we will create an aggregated table in customer level for our transaction data, the goal is to get the clear metric that reflect the behavior and quality of our customer.

=== DATA PREPROCESSING ===

Rows before removing missing customer_ids: 732583

Rows after removing missing customer_ids: 732583

Final dataset shape: (732583, 18)

Date range: 2009-12-01 07:45:00 to 2011-12-09 12:50:00

Analysis period: 738 days

CUSTOMER LEVEL AGGREGATIONS

Customer base size: 5819

Reference date for recency calculation: 2011-12-09 12:50:00

	customer_id	total_orders	first_purchase	last_purchase	\
0	12346.0	2	2010-03-02 13:08:00	2010-06-28 13:53:00	
1	12347.0	8	2010-10-31 14:20:00	2011-12-07 15:52:00	
2	12348.0	5	2010-09-27 14:59:00	2011-09-25 13:13:00	
3	12349.0	3	2010-04-29 13:20:00	2011-11-21 09:51:00	
4	12350.0	1	2011-02-02 16:01:00	2011-02-02 16:01:00	
5	12351.0	1	2010-11-29 15:23:00	2010-11-29 15:23:00	
6	12352.0	7	2010-11-12 10:20:00	2011-11-03 14:37:00	
7	12353.0	2	2010-10-27 12:44:00	2011-05-19 17:47:00	
8	12354.0	1	2011-04-21 13:11:00	2011-04-21 13:11:00	
9	12355.0	2	2010-05-21 11:59:00	2011-05-09 13:49:00	

	total_spent	avg_order_value	total_transactions	total_quantity	\
0	169.36	7.056667	24	24	
1	4921.53	22.169054	222	2967	
2	1658.40	36.052174	46	2704	
3	3405.99	20.895644	163	1435	
4	294.40	18.400000	16	196	
5	300.93	14.330000	21	261	
6	1459.18	18.470633	79	570	
7	406.76	16.948333	24	212	
8	1079.40	18.610345	58	530	
9	947.61	27.074571	35	543	

	unique_products	days_since_first_purchase	days_since_last_purchase	\
0	24	646	528	
1	126	403	1	
2	24	437	74	
3	133	588	18	
4	16	309	309	
5	21	374	374	
6	61	392	35	
7	23	408	203	
8	58	231	231	
9	35	567	213	

	customer_lifespan_days	purchase_frequency
0	118	6.134454
1	402	7.245658
2	362	5.027548
3	570	1.917688
4	0	365.000000
5	0	365.000000
6	356	7.156863
7	204	3.560976
8	0	365.000000
9	353	2.062147

Now we have the aggregated table for our customer, next we will do the RFM (Recency Frequency Monetary) analysis to segment our customer based on their quality



RFM ANALYSIS

RFM Score Distribution:

```
customer_segment
Champions          1703
Lost Customers     1291
Potential Loyalists 969
Loyal Customers    931
At Risk            925
Name: count, dtype: int64
```

	customer_id	recency	frequency	monetary	r_score	f_score	m_score	\
0	12346.0	528	2	169.36	1	2	1	
1	12347.0	1	8	4921.53	5	4	5	
2	12348.0	74	5	1658.40	3	4	4	
3	12349.0	18	3	3405.99	5	3	5	
4	12350.0	309	1	294.40	2	1	2	
5	12351.0	374	1	300.93	2	1	2	
6	12352.0	35	7	1459.18	4	4	4	
7	12353.0	203	2	406.76	2	2	2	
8	12354.0	231	1	1079.40	2	1	3	
9	12355.0	213	2	947.61	2	2	3	

	rfm_score	rfm_score_numeric	customer_segment
0	121	4	Lost Customers
1	545	14	Champions
2	344	11	Loyal Customers
3	535	13	Champions
4	212	5	Lost Customers
5	212	5	Lost Customers
6	444	12	Champions
7	222	6	At Risk
8	213	6	At Risk
9	223	7	At Risk



We have segmented our customer, this reflect a customer quality progression where

1. “Lost Customer” indicates Churned customer
2. “At Risk” have high risk turned into Churned customer in next few months
3. “Potential Customer” their behavior indicates they can be converted into loyal and repeating customer
4. “Loyal Customer” our frequent and repeating customer
5. “Champions” our loyal customer that drives big portion of our revenue

We will use this segmentation for our Customer 360 table

CUSTOMER 360 DATASET CREATION

Final Customer 360 dataset shape: (5819, 22)

Features included: ['customer_id', 'total_orders', 'first_purchase', 'last_purchase', 'total_spent', 'avg_order_value', 'total_transactions', 'total_quantity', 'unique_products', 'days_since_first_purchase', 'days_since_last_purchase', 'customer_lifespan_days', 'purchase_frequency', 'r_score', 'f_score', 'm_score', 'rfm_score', 'rfm_score_numeric', 'customer_segment', 'is_active', 'customer_tenure_months', 'avg_products_per_order']

Customer Segment Distribution:

	Customer_Count	Avg_Revenue	Total_Revenue	Avg_Orders	\
customer_segment					
At Risk	925	517.08	478300.44	1.78	
Champions	1703	6759.63	11511655.36	14.81	
Lost Customers	1291	237.69	306861.37	1.14	
Loyal Customers	931	1517.45	1412747.39	4.65	
Potential Loyalists	969	893.80	866092.69	2.93	

	Avg_Days_Since_Last_Purchase
customer_segment	
At Risk	257.02
Champions	35.18
Lost Customers	456.06
Loyal Customers	109.49
Potential Loyalists	179.77

Sample of Customer 360 Dataset:

	customer_id	total_orders	first_purchase	last_purchase	\
0	12346.0	2	2010-03-02 13:08:00	2010-06-28 13:53:00	
1	12347.0	8	2010-10-31 14:20:00	2011-12-07 15:52:00	
2	12348.0	5	2010-09-27 14:59:00	2011-09-25 13:13:00	
3	12349.0	3	2010-04-29 13:20:00	2011-11-21 09:51:00	
4	12350.0	1	2011-02-02 16:01:00	2011-02-02 16:01:00	
	total_spent	avg_order_value	total_transactions	total_quantity	\
0	169.36	7.056667	24	24	

1	4921.53	22.169054		222	2967
2	1658.40	36.052174		46	2704
3	3405.99	20.895644		163	1435
4	294.40	18.400000		16	196

	unique_products	days_since_first_purchase	...	purchase_frequency	\
0	24		646	...	6.134454
1	126		403	...	7.245658
2	24		437	...	5.027548
3	133		588	...	1.917688
4	16		309	...	365.000000

	r_score	f_score	m_score	rfm_score	rfm_score_numeric	customer_segment	\
0	1	2	1	121	4	Lost Customers	
1	5	4	5	545	14	Champions	
2	3	4	4	344	11	Loyal Customers	
3	5	3	5	535	13	Champions	
4	2	1	2	212	5	Lost Customers	

	is_active	customer_tenure_months	avg_products_per_order
0	False	21.222076	12.000000
1	True	13.239159	15.750000
2	True	14.356110	4.800000
3	True	19.316689	44.333333
4	False	10.151117	16.000000

[5 rows x 22 columns]

4 Phase 2: Customer Segment Deep Dive

Customer segmentation from the RFM analysis is important for us to breakdown the main driver on why customer become churned and how they end up spend more in our product. In this analysis we will characterize each segment and take a look at their behavior

CUSTOMER SEGMENT ANALYSIS ===

Detailed Customer Segment Analysis:

	Count	Avg_Spent	Median_Spent	Total_Spent	Std_Spent	\
customer_segment						
At Risk	925	517.08	410.15	478300.44	540.23	
Champions	1703	6759.63	3365.22	11511655.36	22092.30	
Lost Customers	1291	237.69	204.24	306861.37	159.13	
Loyal Customers	931	1517.45	1190.23	1412747.39	1535.88	
Potential Loyalists	969	893.80	717.21	866092.69	1156.20	

	Avg_Orders	Median_Orders	Std_Orders	Avg_Order_Value	\
customer_segment					
At Risk	1.78	2.0	0.83	33.03	

Champions	14.81	10.0	19.70	32.77
Lost Customers	1.14	1.0	0.35	24.08
Loyal Customers	4.65	4.0	2.73	29.62
Potential Loyalists	2.93	3.0	1.36	29.98

	Median_Order_Value	...	Avg_Recency	Median_Recency	\
customer_segment		...			
At Risk	16.59	...	257.02	240.0	
Champions	18.18	...	35.18	18.0	
Lost Customers	16.08	...	456.06	448.0	
Loyal Customers	16.89	...	109.49	64.0	
Potential Loyalists	17.13	...	179.77	119.0	

	Std_Recency	Avg_Tenure_Months	Median_Tenure_Months	\
customer_segment				
At Risk	185.01	11.52	12.65	
Champions	50.78	18.88	21.29	
Lost Customers	169.41	15.49	15.60	
Loyal Customers	120.37	15.52	17.81	
Potential Loyalists	161.99	13.51	14.49	

	Avg_Purchase_Freq	Median_Purchase_Freq	\
customer_segment			
At Risk	184.07	38.42	
Champions	13.12	7.42	
Lost Customers	328.94	365.00	
Loyal Customers	13.18	4.70	
Potential Loyalists	54.52	6.05	

	Avg_Unique_Products	Median_Unique_Products	Active_Rate
customer_segment			
At Risk	31.65	25.0	0.30
Champions	171.87	130.0	0.91
Lost Customers	18.12	14.0	0.03
Loyal Customers	79.39	63.0	0.62
Potential Loyalists	49.12	39.0	0.45

[5 rows x 21 columns]

SEGMENT IMPACT ANALYSIS

Lost Customers:

- Customer Share: 1,291 customers (22.2%)
- Revenue Share: \$306,861.37 (2.1%)
- Revenue per Customer: \$237.69
- Active Rate: 2.7%
- Avg Recency: 456 days

Champions:

- Customer Share: 1,703 customers (29.3%)
- Revenue Share: \$11,511,655.36 (79.0%)
- Revenue per Customer: \$6759.63
- Active Rate: 91.0%
- Avg Recency: 35 days

Loyal Customers:

- Customer Share: 931 customers (16.0%)
- Revenue Share: \$1,412,747.39 (9.7%)
- Revenue per Customer: \$1517.45
- Active Rate: 61.9%
- Avg Recency: 109 days

At Risk:

- Customer Share: 925 customers (15.9%)
- Revenue Share: \$478,300.44 (3.3%)
- Revenue per Customer: \$517.08
- Active Rate: 29.8%
- Avg Recency: 257 days

Potential Loyalists:

- Customer Share: 969 customers (16.7%)
- Revenue Share: \$866,092.69 (5.9%)
- Revenue per Customer: \$893.80
- Active Rate: 44.6%
- Avg Recency: 180 days

4.0.1 Segment-Level Insights

1. Champions (29.3% of customers, 79% of revenue)

- They are the lifeblood of the business: small in number but extremely high spenders (~\$6,760 each).
- Extremely active (91% active rate, recency ~35 days), with large and frequent orders (15 orders on average, 171 unique products).
- **Risk:** Heavy dependence on this group → losing even a small % will hit revenue hard.

2. Loyal Customers (16% of customers, 9.7% of revenue)

- Mid-value segment, still engaged (62% active, recency ~109 days).
- Spend per customer ~\$1,517, avg. 5 orders, ~79 unique products.
- They are prime candidates to **upgrade into Champions** with personalized offers and upsell campaigns.

3. Potential Loyalists (16.7% of customers, 5.9% of revenue)

- Spend less (~\$894 per customer) and engage moderately (45% active, recency ~180 days).
- Frequency is lower (3 orders, ~49 unique products).

- **Opportunity:** nurture with loyalty incentives, bundles, and cross-sell nudges. They're sitting in the middle, could move upward or churn.

4. At Risk (15.9% of customers, 3.3% of revenue)

- Weak engagement (30% active, recency ~257 days), low spend (~\$517).
- Orders are infrequent (2 orders on average, ~32 unique products).
- They still have some recent touchpoints, so **reactivation campaigns** may save part of this group.

5. Lost Customers (22.2% of customers, only 2.1% of revenue)

- Lowest value segment (~\$238 per customer), inactive for ~456 days.
- Very low activity (1 order, minimal product variety).
- **Impact is minimal on revenue**, so they're not worth aggressive win-back spend. Use low-cost automated reactivation if at all.

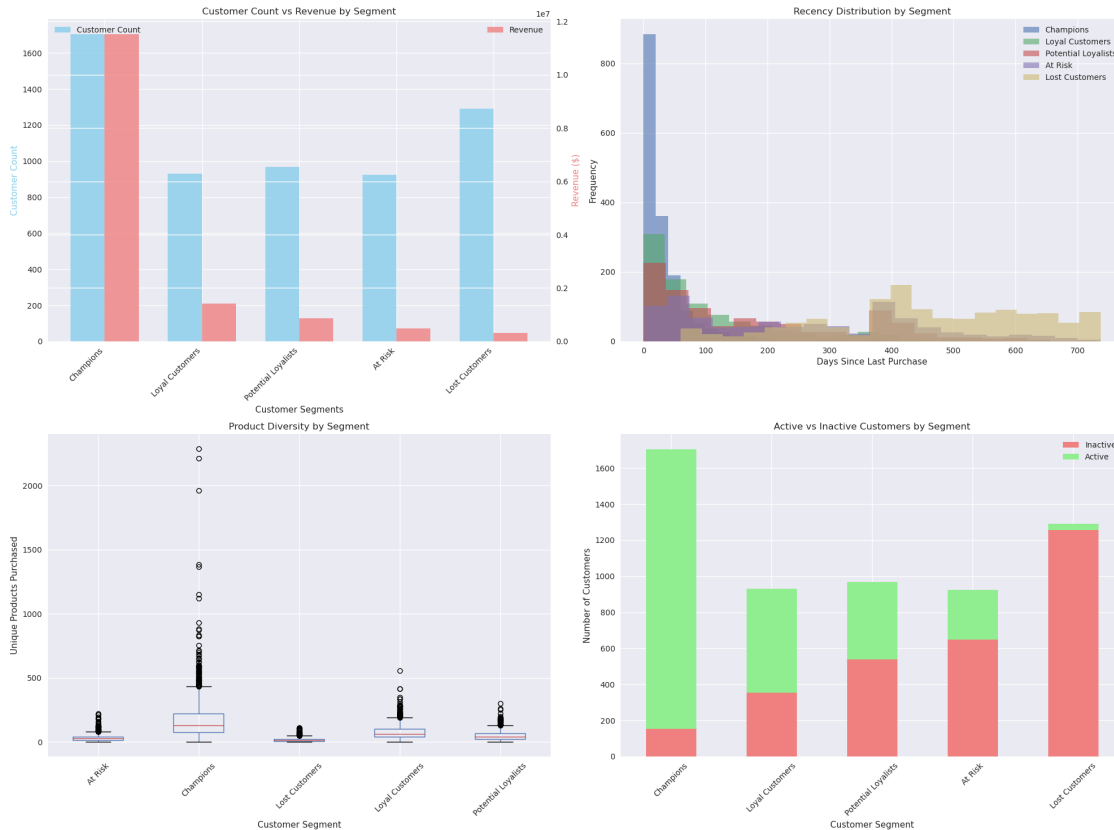
Takeaways

1. **Revenue is highly concentrated:** Champions + Loyal Customers = 45% of customers but 89% of revenue. Business depends heavily on these two groups.
2. **Retention gap is clear:** Potential Loyalists + At Risk represent 32% of the base. If they churn, they'll slide into "Lost" (already 22%). Retention efforts here could protect future revenue.
3. **Lost Customers are low value:** Losing them doesn't hurt much. Focus should remain on retaining and upgrading mid-tier and high-value customers.
4. **Cross-sell opportunity:** Champions buy many unique products (~172), while Potential Loyalists buy ~49. Bridging that gap through tailored product recommendations can accelerate their move upward.
5. **Recency gap:** Champions buy every ~1 month, Loyal every ~3 months, Potential Loyalists ~6 months. Tightening buying cycles through targeted offers could significantly increase revenue.

Recommended Focus:

- **Protect & reward Champions** → VIP perks, early access, loyalty recognition.
- **Upgrade Loyal → Champions** → upselling high-value categories.
- **Nurture Potential Loyalists** → re-engagement campaigns, bundles, cross-sell into sticky products.
- **Selective At-Risk Recovery** → targeted win-back for those with higher past spend.
- **Low-cost automation for Lost** → don't spend much; focus resources upstream.

Do you want me to tie this **Customer 360 segment insight** back to the **Retail Business Analysis & Demand Forecasting** docs you uploaded (like seasonality + product categories)? That way the insight doesn't just stay customer-focused, but also links to **what products and when to target each segment**.



4.1 Segment Purchase Pattern & Product Preference

After we have know the characteristic of each segment, next we want to know if there is any product preference that distinct between these segment and if there is any purchase pattern from the quality customer that can be replicated to lower quality customer

PRODUCT CATEGORY ANALYSIS BY CUSTOMER SEGMENT

Transaction data with segments: 732,583 transactions

Available product categories: 9

Product categories: ['BEAUTY_PERSONAL', 'CHRISTMAS_HOLIDAY', 'FURNITURE_STORAGE', 'GARDEN_OUTDOOR', 'HOME_DECOR', 'KITCHEN_FOOD_UTENSIL', 'STATIONERY_OFFICE', 'TEXTILES_CLOTHING', 'TOYS_GAMES']

CATEGORY PREFERENCES BY SEGMENT

--- Champions Category Preferences ---

Top categories by revenue share:

- HOME_DECOR: 39.8% revenue, 11.9% penetration
- KITCHEN_FOOD_UTENSIL: 20.4% revenue, 11.8% penetration
- TEXTILES_CLOTHING: 9.4% revenue, 11.0% penetration
- BEAUTY_PERSONAL: 8.1% revenue, 11.7% penetration
- GARDEN_OUTDOOR: 6.1% revenue, 11.3% penetration

--- Loyal Customers Category Preferences ---

Top categories by revenue share:

- HOME_DECOR: 38.7% revenue, 12.9% penetration
- KITCHEN_FOOD_UTENSIL: 20.6% revenue, 12.6% penetration
- BEAUTY_PERSONAL: 8.8% revenue, 12.2% penetration
- TEXTILES_CLOTHING: 7.9% revenue, 10.8% penetration
- GARDEN_OUTDOOR: 5.8% revenue, 11.2% penetration

--- Potential Loyalists Category Preferences ---

Top categories by revenue share:

- HOME_DECOR: 37.5% revenue, 14.3% penetration
- KITCHEN_FOOD_UTENSIL: 21.9% revenue, 13.7% penetration
- BEAUTY_PERSONAL: 8.8% revenue, 12.8% penetration
- TEXTILES_CLOTHING: 7.7% revenue, 10.6% penetration
- GARDEN_OUTDOOR: 6.1% revenue, 11.3% penetration

--- At Risk Category Preferences ---

Top categories by revenue share:

- HOME_DECOR: 37.8% revenue, 16.3% penetration
- KITCHEN_FOOD_UTENSIL: 22.8% revenue, 15.0% penetration
- BEAUTY_PERSONAL: 8.6% revenue, 13.3% penetration
- GARDEN_OUTDOOR: 6.4% revenue, 11.1% penetration
- TEXTILES_CLOTHING: 6.2% revenue, 10.1% penetration

--- Lost Customers Category Preferences ---

Top categories by revenue share:

- HOME_DECOR: 38.8% revenue, 19.1% penetration
- KITCHEN_FOOD_UTENSIL: 21.0% revenue, 16.5% penetration
- BEAUTY_PERSONAL: 9.6% revenue, 13.7% penetration
- TEXTILES_CLOTHING: 6.9% revenue, 10.2% penetration
- GARDEN_OUTDOOR: 6.7% revenue, 11.1% penetration

What we learn: Category does not differentiate by loyalty tier; steady contribution but unlikely to drive transitions between segments.

1. Home Decor = Core Engine

- Must always be protected in stock planning, pricing, and promotional visibility.
- Best category to drive initial engagement and repeat purchases across all segments.

2. Kitchen = Wide Appeal but High Churn Risk

- High penetration among At Risk/Lost customers suggests it can't retain on its own.
- Bundle Kitchen products with Décor or Textiles to raise stickiness.

3. Beauty = Entry Point

- Attracts both At Risk and Lost customers.
- Cross-sell Beauty into higher-value categories (e.g., Décor, Kitchen) to prevent attrition.

4. Textiles = Champion Driver

- Disproportionately stronger among Champions, indicating it is part of the upgrade path.
- Push Textiles campaigns at Loyal and Potential Loyalists to encourage category expansion.

Call-to-Action:

- Anchor customer journeys on **Home Decor**, using it as the core of retention campaigns.
- Design **cross-sell flows**
- Focus on **Textiles growth** among mid-tier segments to create Champions.
- Use **low-cost add-ons** (Garden, Beauty) to increase basket size but don't expect them to drive loyalty shifts.

PRODUCT STICKINESS & RETENTION ANALYSIS

1. CATEGORY RETENTION CORRELATION:

Categories ranked by customer activity rate:

	Avg_Recency	Avg_Orders	Avg_Spent	Avg_Lifespan \
product_category				
CHRISTMAS_HOLIDAY	132.82	8.11	3462.13	352.94
FURNITURE_STORAGE	151.12	8.43	3671.05	354.04
TOYS_GAMES	159.30	7.83	3395.62	337.74
TEXTILES_CLOTHING	170.96	7.44	3154.70	324.50
STATIONERY_OFFICE	172.55	7.35	3096.01	322.51
GARDEN_OUTDOOR	175.22	7.29	3075.27	320.49
BEAUTY_PERSONAL	182.48	6.84	2854.14	304.99
KITCHEN_FOOD_UTENSIL	190.57	6.51	2692.57	290.50
HOME_DECOR	195.44	6.26	2577.86	279.90

	Active_Rate
product_category	
CHRISTMAS_HOLIDAY	0.65
FURNITURE_STORAGE	0.60
TOYS_GAMES	0.58
TEXTILES_CLOTHING	0.55
STATIONERY_OFFICE	0.55
GARDEN_OUTDOOR	0.54
BEAUTY_PERSONAL	0.53
KITCHEN_FOOD_UTENSIL	0.51
HOME_DECOR	0.50

2. EARLY PURCHASE CATEGORY IMPACT:

Impact of first purchase category on customer outcomes:

	Customer_Count	Avg_Orders	Avg_CLV	Active_Rate
product_category				
CHRISTMAS_HOLIDAY	277	5.66	2011.12	0.57
STATIONERY_OFFICE	379	5.02	1681.16	0.50
FURNITURE_STORAGE	161	5.63	2052.00	0.50
GARDEN_OUTDOOR	419	6.48	2154.29	0.50
HOME_DECOR	2247	6.14	2553.26	0.50
BEAUTY_PERSONAL	544	6.48	2670.06	0.49

TEXTILES_CLOTHING	471	6.51	2826.32	0.49
KITCHEN_FOOD_UTENSIL	1111	5.91	2325.82	0.48
TOYS_GAMES	210	6.96	4969.11	0.43

3. CHAMPIONS VS LOST CUSTOMERS CATEGORY PREFERENCES:

Categories favored by Champions vs Lost Customers:

Most Champions-favored categories:

- CHRISTMAS_HOLIDAY: Champions 10.6% vs Lost 5.4% (+5.2pp)
- FURNITURE_STORAGE: Champions 10.0% vs Lost 5.6% (+4.4pp)
- TOYS_GAMES: Champions 10.4% vs Lost 8.0% (+2.4pp)
- STATIONERY_OFFICE: Champions 11.3% vs Lost 10.4% (+0.8pp)
- TEXTILES_CLOTHING: Champions 11.0% vs Lost 10.2% (+0.8pp)

Most Lost-Customer-favored categories:

- BEAUTY_PERSONAL: Champions 11.7% vs Lost 13.7% (-2.0pp)
- KITCHEN_FOOD_UTENSIL: Champions 11.8% vs Lost 16.5% (-4.7pp)
- HOME_DECOR: Champions 11.9% vs Lost 19.1% (-7.2pp)

4. CATEGORY CROSS-SELLING PATTERNS:

Impact of category diversity on customer value:

	Customer_Count	Avg_CLV	Avg_Orders	Active_Rate
categories_purchased				
1	182	522.30	1.71	0.21
2	189	457.23	1.74	0.25
3	320	509.38	1.88	0.29
4	362	556.81	2.10	0.30
5	495	881.15	2.48	0.28
6	609	951.73	2.89	0.36
7	765	1176.24	3.85	0.40
8	999	1906.84	5.55	0.49
9	1898	5378.87	11.60	0.75

Optimal category diversity for CLV: 9 categories

CLV at optimal diversity: \$5378.87

5. SEQUENTIAL PURCHASE PATTERNS:

Same-category repeat purchase rate: 40.7%

Top category transitions (First → Second purchase):

- HOME_DECOR → HOME_DECOR: 196 customers
- KITCHEN_FOOD_UTENSIL → KITCHEN_FOOD_UTENSIL: 81 customers
- HOME_DECOR → KITCHEN_FOOD_UTENSIL: 62 customers
- KITCHEN_FOOD_UTENSIL → HOME_DECOR: 57 customers
- TEXTILES_CLOTHING → TEXTILES_CLOTHING: 41 customers
- HOME_DECOR → TEXTILES_CLOTHING: 32 customers
- HOME_DECOR → BEAUTY_PERSONAL: 28 customers
- BEAUTY_PERSONAL → HOME_DECOR: 27 customers

- HOME_DECOR → GARDEN_OUTDOOR: 24 customers
- GARDEN_OUTDOOR → HOME_DECOR: 22 customers

4.1.1 1. Category Retention & Stickiness

- **High-Activity Categories** (best for retention):
 - **Christmas/Holiday, Furniture/Storage, Toys/Games** → Customers who buy here show higher order counts (~8+) and longer lifespans (~350 days).
 - These categories **correlate with “Champions”** – they bring stronger ongoing activity.
- **Weak Stickiness Categories:**
 - **Home Decor & Kitchen/Food** dominate revenue, but customers here are **less active** (avg. 6 orders, lifespan ~280 days).
 - They are big **acquisition funnels** but underperform in converting to loyal customers.

4.1.2 2. First Purchase Impact

- **High-Value First Purchases:**
 - Customers starting in **Toys/Games** have the **highest CLV** (\$4,969) but poor retention (active rate 0.43).
 - **Kitchen/Food & Beauty/Personal** → balance between volume and stickiness; good entry points for loyalty.
- **Lower Value Onboarding:**
 - **Stationery & Garden** yield lower CLV (<\$2,200) despite decent activity.

4.1.3 3. Champions vs Lost Customers

- **Champions** buys seasonal/niche: Christmas, Furniture, Toys.
- **Lost Customers** skew toward everyday staples: Home Décor, Kitchen, Beauty.

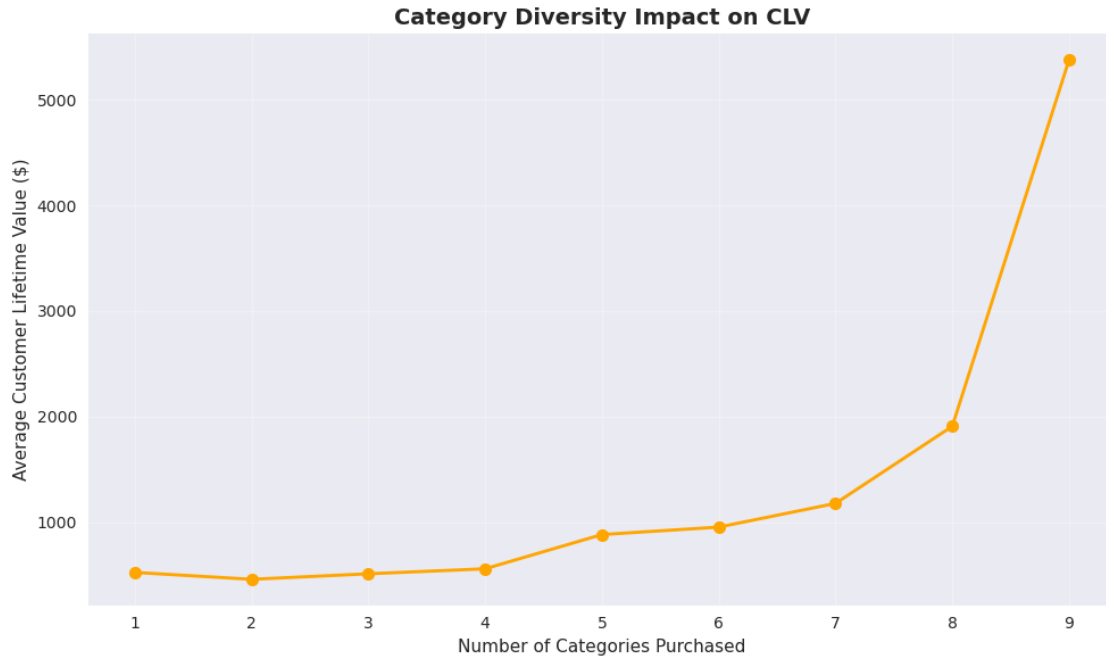
4.1.4 4. Cross-Selling & Category Diversity

- Strong positive relationship:
 - CLV grows dramatically with category diversity.
 - Customers buying from **9 categories** are worth **~\$5,379 CLV**, 10× higher than single-category buyers.
- Active rate rises from **21% (1 category)** → **75% (9 categories)**.

Cross-category engagement is the single strongest lever for retention and value. Campaigns must actively push multi-category shopping.


4.1.5 5. Sequential Purchase Patterns

- **Same-category repeat rate = 41%** → customers like to re-buy within the same category (esp. Home Decor, Kitchen).



Here we can see from the chart that Customer Lifetime Value (CLV) grows almost exponentially as customers purchase across more categories. For single-category buyers, CLV is very low (~\$500), and it stays relatively flat up to 4 categories. But starting from 5+ categories, CLV rises steeply, reaching nearly \$5,400 when customers buy across all 9 categories.

Customers who diversify their purchases are not just buying more often, but also staying active longer, which compounds their lifetime value.



Everyday categories (Home Decor, Kitchen, Beauty) are essential for acquisition but insufficient for retention. Without cross-sell into higher-stickiness categories, these customers churn.

Seasonal/niche categories (Christmas, Furniture, Toys) are associated with Champions and should be used as “retention hooks.”

5 Phase 3: Machine Learning for Churn Prediction

In this section, we’ll develop and evaluate machine learning models to predict customer churn

1. **Data Preparation & Feature Engineering** - Create predictive features without data leakage
2. **Data Splitting Strategy** - Split the data, predict the last 90 days of data
3. **Model Training** - Train multiple algorithms (Logistic Regression, Random Forest, XGBoost)
4. **Model Evaluation & Comparison** - Compare performance and select best model
5. **Final Predictions** - Generate churn scores and recommendations

5.1 Data Preparation & Feature Engineering

5.1.1 Import Required Libraries

5.1.2 Temporal Data Split (Preventing Data Leakage)

Data Timeline:

- Full period: 2009-12-01 to 2011-12-09 (738 days)
- Training period: 2009-12-01 to 2011-07-14
- Validation period: 2011-07-14 to 2011-12-09
- Prediction window: 90 days
- Active threshold: 90 days (exclude already churned customers)

Training Data (Excluding Already Churned):

- All training transactions: 531,153
- All customers in training period: 5,022
- Active customers at split date: 1,923
- Already churned (excluded): 3,099
- Active training transactions: 347,854

Churn Target Definition (Active Customers Only):

- Active customers at split: 1,923
- Customers active after split: 1,436
- Customers who churned after split: 487
- Churn rate (among active): 25.3%

5.1.3 Feature Engineering

Create predictive features using only data available up to the split date.

Feature Engineering Complete:

- Total customers: 1,923
- Total features: 24
- Churn rate: 25.3%

Feature Categories:

- Volume features: total_orders, total_transactions, total_quantity
- Monetary features: total_spent, avg_order_value, monetary_per_day
- Behavioral features: purchase_frequency, product_diversity,

category_diversity

- Temporal features: days_since_first, days_since_last, customer_lifespan
- Early behavior: early_revenue, early_orders, early_products,

early_categories

	customer_id	total_orders	first_purchase	last_purchase_in_training	\
0	12347.0	5	2010-10-31 14:20:00	2011-06-09 13:01:00	
1	12353.0	2	2010-10-27 12:44:00	2011-05-19 17:47:00	
2	12354.0	1	2011-04-21 13:11:00	2011-04-21 13:11:00	
3	12355.0	2	2010-05-21 11:59:00	2011-05-09 13:49:00	
4	12358.0	4	2009-12-08 07:59:00	2011-07-12 10:04:00	

	total_transactions	total_spent	avg_order_value	spend_volatility	\
0	142	2817.48	19.841408	21.527576	
1	24	406.76	16.948333	8.799589	
2	58	1079.40	18.610345	8.679742	
3	35	947.61	27.074571	22.110475	
4	68	2923.87	42.998088	64.174167	

	total_quantity	avg_quantity_per_order	...	monetary_per_day	\
0	1822	12.830986	...	11.005781	
1	212	8.833333	...	1.564462	
2	530	9.137931	...	12.850000	
3	543	15.514286	...	2.261599	
4	924	13.588235	...	5.015214	

	product_diversity	category_diversity	avg_days_between_orders	\
0	16.833333	0.9	36.666667	
1	7.666667	0.5	68.000000	
2	29.000000	0.8	0.000000	
3	11.666667	0.7	117.666667	
4	11.000000	0.6	116.200000	

	early_revenue	early_orders	early_products	early_categories	\
0	611.53	1	40	7	
1	317.76	1	20	5	
2	1079.40	1	58	8	

3	488.21	1	22	6
4	1429.83	1	17	5

	early_avg_order_value	churned
0	15.288250	0
1	15.888000	1
2	18.610345	1
3	22.191364	1
4	84.107647	0

[5 rows x 26 columns]

5.2 Train Test Data Splitting Strategy

Dataset Prepared:

- Feature matrix shape: (1923, 22)
- Target vector shape: (1923,)
- Features: 22
- Positive class (churned): 487 (25.3%)
- Negative class (active): 1,436 (74.7%)

Train-Validation Split:

- Training set: 1,538 customers (80.0%)
- Validation set: 385 customers (20.0%)
- Training churn rate: 25.3%
- Validation churn rate: 25.5%

Data ready for model training

- X_train: (1538, 22) (original)
- X_train_scaled: (1538, 22) (standardized)
- X_val: (385, 22) (original)
- X_val_scaled: (385, 22) (standardized)

5.3 Model Training

Train and tune three different machine learning algorithms to predict customer churn.

5.3.1 Logistic Regression

Logistic Regression Results:

- AUC-ROC: 0.786
- Average Precision: 0.549
- F1-Score: 0.542
- CV AUC: 0.769 ± 0.029

Top 5 Most Important Features (Logistic Regression):

- | | |
|-----------------------|-------------------------------------|
| 6. total_quantity | Decreases churn risk (coef: -1.515) |
| 2. total_transactions | Decreases churn risk (coef: -1.144) |

3. total_spent	Decreases churn risk (coef: -0.995)
17. avg_days_between_orders	Increases churn risk (coef: +0.677)
8. unique_products	Increases churn risk (coef: +0.618)

Logistic Regression trained successfully

Logistic Regression was chosen first because it's **simple, interpretable, and fast**. It sets a clear baseline before moving to more complex models (like Random Forest or XGBoost) which may improve accuracy but sacrifice interpretability.

1. Class Imbalance

- Churn datasets are often imbalanced (fewer churners than non-churners).
- We used `class_weight='balanced'` so the model pays equal attention to both classes, avoiding bias toward the majority.

2. Feature Scaling

- Logistic Regression is sensitive to feature scale.
- Input features were standardized (`X_train_scaled`) to ensure coefficients are comparable and model converges efficiently.

3. Interpretability

- One strength of Logistic Regression is **transparent coefficients**.
- Positive coefficients → increase churn risk; negative → reduce churn risk. This helps translate results into business insights (e.g., more orders reduce churn).

4. Evaluation Metrics

- Used multiple metrics to balance perspectives:
 - **AUC-ROC (0.786)**: ability to rank churn vs non-churn.
 - **Average Precision (0.549)**: robustness in imbalanced setting.
 - **F1-score (0.542)**: balance of precision and recall.
- Cross-validation confirmed stability (AUC $\sim 0.769 \pm 0.029$).

5.3.2 Random Forest

Random Forest Results:

- AUC-ROC: 0.765
- Average Precision: 0.514
- F1-Score: 0.525
- CV AUC: 0.772 ± 0.030

Top 5 Most Important Features (Random Forest):

3. total_spent	(importance: 0.136)
6. total_quantity	(importance: 0.117)
1. total_orders	(importance: 0.083)
14. monetary_per_day	(importance: 0.081)
2. total_transactions	(importance: 0.076)

Random Forest trained successfully

Unlike Logistic Regression, Random Forest handles raw feature scales directly. More flexible in capturing **nonlinear patterns** and feature interactions.

1. Class Imbalance

- `class_weight='balanced'` was applied to prevent bias toward the majority (non-churners).

2. Hyperparameters

- `max_depth=10`, `min_samples_split=20`, `min_samples_leaf=10` → constraints added to **avoid overfitting** while keeping interpretability of feature importance.
- `n_estimators=100` ensures stability of predictions.

3. Performance

- **AUC-ROC = 0.765, AP = 0.514, F1 = 0.525.**
- Slightly weaker than Logistic Regression (AUC 0.786), but **cross-validation AUC $\sim 0.772 \pm 0.030$** shows stable generalization.

5.3.3 XGBoost

XGBoost Results:

- AUC-ROC: 0.739
- Average Precision: 0.465
- F1-Score: 0.483
- CV AUC: 0.739 ± 0.019

Top 5 Most Important Features (XGBoost):

- | | |
|-----------------------------|---------------------|
| 3. total_spent | (importance: 0.114) |
| 6. total_quantity | (importance: 0.080) |
| 17. avg_days_between_orders | (importance: 0.060) |
| 1. total_orders | (importance: 0.057) |
| 13. purchase_frequency | (importance: 0.049) |

XGBoost trained successfully

XGBoost builds trees sequentially, correcting mistakes from previous trees. This allows it to capture nonlinear patterns and interactions much better than Logistic Regression or Random Forest.

1. Class Imbalance

- Used `scale_pos_weight = (# non-churners / # churners)` to balance churn prediction, ensuring the model does not just predict the majority class.

2. Regularization & Stability

- Parameters (`max_depth=6`, `subsample=0.8`, `colsample_bytree=0.8`) help reduce overfitting.
- `learning_rate=0.1` controls how fast the model adapts, trading off speed vs generalization.

Overall the predictive power is slightly weaker compared to Logistic Regression (AUC ~0.786) and roughly on par with Random Forest. A caveat is that XGBoost is more data-hungry and parameter-sensitive. With richer behavioral or marketing features (campaign response, engagement signals), it may show stronger advantages. For now, Logistic Regression remains the best balance of accuracy and interpretability.

5.4 Model Evaluation and Comparison

	Model	Recall	Precision	F1-Score	AUC-ROC	\
0	Logistic Regression	0.755102	0.422857	0.542125	0.785999	
1	Random Forest	0.653061	0.438356	0.524590	0.765022	
2	XGBoost	0.500000	0.466667	0.482759	0.739245	

	Average Precision
0	0.548980
1	0.513514
2	0.465052

BEST MODEL: Logistic Regression

- Best Recall: 0.755 (75.5% of churners detected)
- AUC-ROC: 0.786
- Model Type: Linear
- Selection Criteria: Highest recall for churn detection

DETAILED EVALUATION - Logistic Regression:

Classification Report:

	precision	recall	f1-score	support
Active	0.89	0.65	0.75	287
Churned	0.42	0.76	0.54	98
accuracy			0.68	385
macro avg	0.65	0.70	0.65	385
weighted avg	0.77	0.68	0.70	385

Confusion Matrix:

	Predicted	
Actual	Active	Churned
Active	186	101
Churned	24	74

Business Metrics:

- RECALL (Sensitivity): 75.5% of churners identified
- Precision: 42.3% of churn predictions are correct
- Specificity: 64.8% of active customers correctly identified

PREDICTING CHURN FOR NEXT 3 MONTHS

Analysis date: 2011-12-09

Predicting churn for 3-month period: 2011-12-09 to 2012-03-09

Creating features for prediction (matching training features)...

Calculating early behavior features...

Features created for prediction: ['total_orders', 'total_transactions', 'total_spent', 'avg_order_value', 'spend_volatility', 'total_quantity', 'avg_quantity_per_order', 'unique_products', 'unique_categories', 'days_since_first', 'days_since_last', 'customer_lifespan', 'purchase_frequency', 'monetary_per_day', 'product_diversity', 'category_diversity', 'avg_days_between_orders', 'early_revenue', 'early_orders', 'early_products', 'early_categories', 'early_avg_order_value']

Shape of prediction dataset: (5819, 22)

Excluding 'Lost Customers' from churn prediction list...

Customers before excluding 'Lost Customers': 5,819

Customers after excluding 'Lost Customers': 4,528

CHURN PREDICTION SUMMARY (Next 3 Months) - HIGH RISK ONLY (>80%):

Total customers analyzed (excluding Lost Customers): 4,528

High risk customers (>80%): 1,556

High risk rate: 34.4%

Risk Distribution (High Risk Only):

- High Risk: 527 customers (33.9%)
- Critical Risk: 1,029 customers (66.1%)

Segment-wise High Risk Distribution:

	Avg_Churn_Prob	Customer_Count
customer_segment		
At Risk	0.942	690
Potential Loyalists	0.918	599
Loyal Customers	0.884	254
Champions	0.878	13

High-Value Customers at High Risk (Top 20):

	customer_id	customer_segment	churn_probability	total_spent	\
4412	18139.0	Champions	0.999134	8438.34	
2727	16000.0	Champions	0.999069	12393.70	
1944	14938.0	At Risk	0.998777	1757.31	
3269	16716.0	At Risk	0.998428	1248.48	
3744	17305.0	At Risk	0.997938	2135.46	
2605	15823.0	Potential Loyalists	0.997655	3217.21	
315	12742.0	At Risk	0.996891	1185.02	
3539	17039.0	At Risk	0.996403	1954.99	
1864	14831.0	Potential Loyalists	0.996378	1440.12	

2538	15736.0	At Risk	0.995937	1682.17
2813	16118.0	At Risk	0.995643	3997.73
260	12671.0	At Risk	0.994823	2622.48
650	13205.0	At Risk	0.994502	2803.20
2295	15413.0	Loyal Customers	0.994470	6798.72
1956	14956.0	At Risk	0.994204	1325.00
910	13543.0	At Risk	0.994127	1439.61
4348	18051.0	Potential Loyalists	0.994015	1863.48
3864	17448.0	Loyal Customers	0.993846	13928.02
437	12911.0	At Risk	0.993644	1651.72
747	13337.0	At Risk	0.993296	1550.06

	total_orders	days_since_last	risk_category
4412	6	17	Critical Risk
2727	3	2	Critical Risk
1944	2	561	Critical Risk
3269	1	630	Critical Risk
3744	1	648	Critical Risk
2605	2	728	Critical Risk
315	1	625	Critical Risk
3539	1	603	Critical Risk
1864	3	679	Critical Risk
2538	2	434	Critical Risk
2813	1	652	Critical Risk
260	1	605	Critical Risk
650	1	442	Critical Risk
2295	5	691	Critical Risk
1956	1	420	Critical Risk
910	2	633	Critical Risk
4348	7	633	Critical Risk
3864	41	528	Critical Risk
437	1	551	Critical Risk
747	1	545	Critical Risk

Total revenue at risk from high-value customers: \$73,432.82

High Risk Customers Summary:

Count: 1,556 customers

Total spent at risk: \$1,522,389.17

Average spend per at-risk customer: \$978.40

High risk churn prediction complete. Results saved for 1,556 customers.

Here we can see from the model comparison that Logistic Regression gave the best balance of accuracy (AUC ~0.786) and interpretability, while Random Forest (AUC ~0.765) and XGBoost (AUC ~0.739) provided additional validation but did not significantly outperform. All three models highlight the same key churn drivers: higher spend, quantity, and transaction frequency reduce churn, while longer purchase gaps increase churn risk.

The insight here is that churn in this business is strongly linked to engagement depth and consistency. Customers who are active across multiple purchases and categories remain sticky, while those who disengage for long periods are at high risk of churn.

What we learn is that our models are reliable enough to flag customers at risk, but precision is not perfect. This creates an opportunity to over-predict churn deliberately: treating a broader group of customers as “at risk” may waste some effort on false positives, but it ensures we don’t miss truly at-risk customers.

The next step is to use churn scores to over-identify at-risk customers, then reactivate them with sticky-category recommendations.

6 Phase 4: Recommendation System

In this phase, we design a simple recommendation system aligned with three key business scenarios, ensuring recommendations support revenue growth, customer stickiness, and churn reactivation.

We target different segment with different scenario

- * Scenario 1: Champions and Loyal Customer (Drive Revenue)
- * Scenario 2: At Risk and Potential Loyal (Stickiness)
- * Scenario 3: Churn Prevention

Scenario 1: Drive Revenue

- **Logic:**
 1. Identify customer’s **favorite category** (highest spending).
 2. Retrieve **top-selling products** in that category.
 3. Recommend **3 products the customer has not purchased**.
- **Goal:** Increase revenue by deepening spend in customer’s preferred area.

Scenario 2: Cross-Selling (Stickiness)

- **Logic:**
 1. Start from customer’s **favorite category**.
 2. Identify **related categories** (via correlation analysis).
 3. Recommend **1 top product from each related category** (3 in total).
- **Goal:** Broaden category engagement, driving **higher stickiness and CLV**.

Scenario 3: Churn Prevention

- **Logic:**
 1. Use churn model to identify **high-risk customers**.
 2. Focus on **high-stickiness categories** (>60% retention).
 3. Recommend **Top 3 products they haven’t tried** from these categories.
 4. Offer promo to win back high risk customers
- **Goal:** Win back disengaged customers with proven retention drivers.

And we will also consider seasonal item * Exclude **Christmas/Holiday** products for recommendations outside the festive period. * Focus on **year-round categories** for sustainable engagement.

6.0.1 Data Preparation for Recommendation System

First, let's analyze category correlations and prepare the data needed for our recommendation scenarios.

=== RECOMMENDATION SYSTEM DATA PREPARATION ===

1. CATEGORY CORRELATION ANALYSIS:

Category correlation matrix:

product_category	BEAUTY_PERSONAL	CHRISTMAS_HOLIDAY	FURNITURE_STORAGE	\
product_category				
BEAUTY_PERSONAL	1.000	0.668	0.803	
CHRISTMAS_HOLIDAY	0.668	1.000	0.528	
FURNITURE_STORAGE	0.803	0.528	1.000	
GARDEN_OUTDOOR	0.806	0.425	0.693	
HOME_DECOR	0.874	0.501	0.791	
KITCHEN_FOOD_UTENSIL	0.875	0.704	0.795	
STATIONERY_OFFICE	0.779	0.598	0.724	
TEXTILES_CLOTHING	0.633	0.595	0.514	
TOYS_GAMES	0.657	0.667	0.551	

product_category	GARDEN_OUTDOOR	HOME_DECOR	KITCHEN_FOOD_UTENSIL	\
product_category				
BEAUTY_PERSONAL	0.806	0.874	0.875	
CHRISTMAS_HOLIDAY	0.425	0.501	0.704	
FURNITURE_STORAGE	0.693	0.791	0.795	
GARDEN_OUTDOOR	1.000	0.960	0.723	
HOME_DECOR	0.960	1.000	0.820	
KITCHEN_FOOD_UTENSIL	0.723	0.820	1.000	
STATIONERY_OFFICE	0.675	0.761	0.806	
TEXTILES_CLOTHING	0.446	0.544	0.708	
TOYS_GAMES	0.467	0.565	0.757	

product_category	STATIONERY_OFFICE	TEXTILES_CLOTHING	TOYS_GAMES
product_category			
BEAUTY_PERSONAL	0.779	0.633	0.657
CHRISTMAS_HOLIDAY	0.598	0.595	0.667
FURNITURE_STORAGE	0.724	0.514	0.551
GARDEN_OUTDOOR	0.675	0.446	0.467
HOME_DECOR	0.761	0.544	0.565
KITCHEN_FOOD_UTENSIL	0.806	0.708	0.757
STATIONERY_OFFICE	1.000	0.605	0.612
TEXTILES_CLOTHING	0.605	1.000	0.762
TOYS_GAMES	0.612	0.762	1.000

2. CATEGORY STICKINESS ANALYSIS:

Category stickiness (retention rate):

	category	retention_rate
0	CHRISTMAS_HOLIDAY	0.654719
5	FURNITURE_STORAGE	0.595630
7	TOYS_GAMES	0.576903
8	TEXTILES_CLOTHING	0.554923
6	STATIONERY_OFFICE	0.546602
4	GARDEN_OUTDOOR	0.540821
1	BEAUTY_PERSONAL	0.529591
3	KITCHEN_FOOD_UTENSIL	0.509516
2	HOME_DECOR	0.502674

High-stickiness categories (>60% retention): ['CHRISTMAS_HOLIDAY']

3. PRODUCT PERFORMANCE ANALYSIS:

Product performance calculated for 4755 products

Top 5 products by popularity score:

	product_id	product_category	total_revenue	unique_customers	\
4215	85123A	HOME_DECOR	148184.57	1395	
1652	22423	KITCHEN_FOOD_UTENSIL	125625.55	1121	
3938	84879	HOME_DECOR	121829.55	995	
4193	85099B	TEXTILES_CLOTHING	119378.46	930	
3163	47566	KITCHEN_FOOD_UTENSIL	95829.48	874	

	popularity_score
4215	1.000000
1652	0.830092
3938	0.778593
4193	0.750031
3163	0.638623

4. CUSTOMER PURCHASE HISTORY:

Purchase history compiled for 5819 customers

5. SEASONAL PRODUCT IDENTIFICATION:

Seasonal categories to filter: ['CHRISTMAS_HOLIDAY']

Current month: 12, Is festive period: True

6.0.2 Scenario 1: Drive Revenue (Champions and Loyal Customers)

Logic: 1. Identify customer's favorite category (highest spending) 2. Retrieve top-selling products in that category 3. Recommend 3 products the customer has not purchased

Goal: Increase revenue by deepening spend in customer's preferred area.

=== SCENARIO 1: DRIVE REVENUE RECOMMENDATIONS ===

Target: Champions and Loyal Customers

Generating recommendations for 2634 Champions and Loyal Customers...

Scenario 1 Results:

- Customers targeted: 2,634
- Recommendations generated: 7,902
- Average recommendations per customer: 3.0
- Categories recommended: 9

Top categories in Scenario 1 recommendations:

- HOME_DECOR: 6231 recommendations
- KITCHEN_FOOD_UTENSIL: 987 recommendations
- TEXTILES_CLOTHING: 345 recommendations
- CHRISTMAS_HOLIDAY: 117 recommendations
- STATIONERY_OFFICE: 66 recommendations

Sample Scenario 1 Recommendations:

	customer_id	customer_segment	recommended_product_id	recommended_category	\
0	12347.0	Champions	85123A	HOME_DECOR	
1	12347.0	Champions	84879	HOME_DECOR	
2	12347.0	Champions	22469	HOME_DECOR	
3	12348.0	Loyal Customers	22423	KITCHEN_FOOD_UTENSIL	
4	12348.0	Loyal Customers	47566	KITCHEN_FOOD_UTENSIL	
5	12348.0	Loyal Customers	21212	KITCHEN_FOOD_UTENSIL	
6	12349.0	Champions	85123A	HOME_DECOR	
7	12349.0	Champions	22469	HOME_DECOR	
8	12349.0	Champions	22138	HOME_DECOR	
9	12352.0	Champions	85123A	HOME_DECOR	

	reason
0	Top product in favorite category (HOME_DECOR)
1	Top product in favorite category (HOME_DECOR)
2	Top product in favorite category (HOME_DECOR)
3	Top product in favorite category (KITCHEN_FOOD...
4	Top product in favorite category (KITCHEN_FOOD...
5	Top product in favorite category (KITCHEN_FOOD...
6	Top product in favorite category (HOME_DECOR)
7	Top product in favorite category (HOME_DECOR)
8	Top product in favorite category (HOME_DECOR)
9	Top product in favorite category (HOME_DECOR)

Scenario 1 implementation completed!

6.0.3 Scenario 2: Cross-Selling for Stickiness (At Risk and Potential Loyalists)

Logic: 1. Start from customer's favorite category 2. Identify related categories (via correlation analysis) 3. Recommend 1 top product from each related category (3 in total)

Goal: Broaden category engagement, driving higher stickiness and CLV.

=== SCENARIO 2: CROSS-SELLING FOR STICKINESS RECOMMENDATIONS ===

Target: At Risk and Potential Loyalists

Generating recommendations for 1894 At Risk and Potential Loyalists...

Scenario 2 Results:

- Customers targeted: 1,894
- Recommendations generated: 2,938
- Average recommendations per customer: 1.6
- Categories recommended: 9

Top categories in Scenario 2 recommendations:

- FURNITURE_STORAGE: 821 recommendations
- TOYS_GAMES: 594 recommendations
- STATIONERY_OFFICE: 465 recommendations
- GARDEN_OUTDOOR: 393 recommendations
- BEAUTY_PERSONAL: 326 recommendations

Recommendation reasons breakdown:

- New category recommendations: 2903
- High retention category recommendations: 35

Sample Scenario 2 Recommendations:

	customer_id	customer_segment	recommended_product_id	recommended_category	\
0	12353.0	At Risk	21754	FURNITURE_STORAGE	
1	12353.0	At Risk	21791	TOYS_GAMES	
2	12353.0	At Risk	23298	GARDEN_OUTDOOR	
3	12355.0	At Risk	21754	FURNITURE_STORAGE	
4	12361.0	At Risk	21755	BEAUTY_PERSONAL	
5	12361.0	At Risk	48138	STATIONERY_OFFICE	
6	12361.0	At Risk	22086	CHRISTMAS_HOLIDAY	
7	12363.0	At Risk	21755	BEAUTY_PERSONAL	
8	12363.0	At Risk	21791	TOYS_GAMES	
9	12363.0	At Risk	85099B	TEXTILES_CLOTHING	

	reason
0	Cross-sell from KITCHEN_FOOD_UTENSIL to FURNIT...
1	Cross-sell from KITCHEN_FOOD_UTENSIL to TOYS_G...
2	Cross-sell from KITCHEN_FOOD_UTENSIL to GARDEN...
3	Cross-sell from KITCHEN_FOOD_UTENSIL to FURNIT...
4	Cross-sell from TEXTILES_CLOTHING to BEAUTY_PE...
5	Cross-sell from TEXTILES_CLOTHING to STATIONER...
6	Cross-sell from TEXTILES_CLOTHING to CHRISTMAS...
7	Cross-sell from CHRISTMAS_HOLIDAY to BEAUTY_PE...
8	Cross-sell from CHRISTMAS_HOLIDAY to TOYS_GAME...
9	Cross-sell from CHRISTMAS_HOLIDAY to TEXTILES_...

Scenario 2 implementation completed!

6.0.4 Scenario 3: Churn Prevention (High-Risk Customers)

Logic: 1. Use churn model to identify high-risk customers 2. Focus on high-stickiness categories (>60% retention) 3. Recommend top 3 products they haven't tried from these categories 4. Offer promotional incentives to win back high risk customers

Goal: Win back disengaged customers with proven retention drivers.

=== SCENARIO 3: CHURN PREVENTION RECOMMENDATIONS ===

Target: Customers Predicted to Churn (90-day prediction)

1. USING EXISTING CHURN PREDICTIONS:

- Using customers from next_3_month_churn_list
- Customers predicted to churn (90-day window): 1556
- Average churn probability: 0.923

2. GENERATING CHURN PREVENTION RECOMMENDATIONS:

Targeting 1556 customers predicted to churn (90-day window)

Scenario 3 Results:

- Customers targeted: 1,556
- Recommendations generated: 4,668
- Average recommendations per customer: 3.0
- Categories recommended: 5

Top categories in Scenario 3 recommendations:

- FURNITURE_STORAGE: 1497 recommendations
- TOYS_GAMES: 1446 recommendations
- TEXTILES_CLOTHING: 1277 recommendations
- STATIONERY_OFFICE: 265 recommendations
- GARDEN_OUTDOOR: 183 recommendations

Sample Scenario 3 Recommendations:

	customer_id	recommended_product_id	recommended_category	churn_probability	\
0	17945.0	21754	FURNITURE_STORAGE	0.999939	
1	17945.0	21791	TOYS_GAMES	0.999939	
2	17945.0	85099B	TEXTILES_CLOTHING	0.999939	
3	15959.0	21754	FURNITURE_STORAGE	0.999933	
4	15959.0	21791	TOYS_GAMES	0.999933	
5	15959.0	85099B	TEXTILES_CLOTHING	0.999933	
6	15794.0	21791	TOYS_GAMES	0.999817	
7	15794.0	23298	GARDEN_OUTDOOR	0.999817	
8	15794.0	21754	FURNITURE_STORAGE	0.999817	
9	13526.0	21754	FURNITURE_STORAGE	0.999781	
	risk_level	promotional_offer			
0	High	25% discount + free shipping			
1	High	25% discount + free shipping			

2	High	25% discount + free shipping
3	High	25% discount + free shipping
4	High	25% discount + free shipping
5	High	25% discount + free shipping
6	High	25% discount + free shipping
7	High	25% discount + free shipping
8	High	25% discount + free shipping
9	High	25% discount + free shipping

Scenario 3 implementation completed!

6.0.5 Final Customer 360 Recommendation Table

Now let's create the final comprehensive table that combines all scenarios and shows each customer with their segment and top 3 product recommendations based on their appropriate scenario.

=== FINAL CUSTOMER 360 RECOMMENDATION TABLE ===

1. COMBINING ALL RECOMMENDATION SCENARIOS:

- Scenario 1 (Revenue): 7902 recommendations
- Scenario 2 (Stickiness): 2938 recommendations
- Scenario 3 (Churn Prevention): 4668 recommendations
- Total recommendations: 15508

1.5. CREATING PRODUCT LOOKUP TABLE:

- Product lookup table created with 4604 products

2. CREATING CUSTOMER-LEVEL SUMMARY:

3. RECOMMENDATION SUMMARY STATISTICS:

Customers by primary scenario (hierarchy applied):

- Revenue Growth: 2367 customers
- Churn Prevention: 1556 customers
- No Action: 1291 customers
- Cross-Selling: 474 customers
- No Recommendations: 131 customers

Customers by segment:

- Champions: 1703 customers
- Lost Customers: 1291 customers
- Potential Loyalists: 969 customers
- Loyal Customers: 931 customers
- At Risk: 925 customers

Recommendation coverage:

- Customers with recommendations: 4397 (75.6%)
- Customers without recommendations: 1422
- Average recommendations per customer: 2.9

4. DATA INTEGRITY CHECK:

- Duplicate customers: 0 (should be 0)
- Unique customers: 5819
- Total rows: 5819

5. SAMPLE OF FINAL CUSTOMER 360 RECOMMENDATION TABLE:

Sample customers with recommendations:

	customer_id	customer_segment	total_spent	total_orders	primary_scenario	\
1	12347.0	Champions	4921.53	8	Revenue Growth	
2	12348.0	Loyal Customers	1658.40	5	Revenue Growth	
3	12349.0	Champions	3405.99	3	Revenue Growth	
6	12352.0	Champions	1459.18	7	Revenue Growth	
7	12353.0	At Risk	406.76	2	Churn Prevention	
8	12354.0	At Risk	1079.40	1	Churn Prevention	
9	12355.0	At Risk	947.61	2	Churn Prevention	
10	12356.0	Champions	5611.73	6	Revenue Growth	
11	12357.0	Loyal Customers	17437.66	2	Revenue Growth	
12	12358.0	Champions	3447.07	5	Revenue Growth	

	recommendation_1_product	recommendation_1_name	\
1	85123A	WHITE HANGING HEART T-LIGHT HOLDER	
2	22423	REGENCY CAKESTAND 3 TIER	
3	85123A	WHITE HANGING HEART T-LIGHT HOLDER	
6	85123A	WHITE HANGING HEART T-LIGHT HOLDER	
7	21754	HOME BUILDING BLOCK WORD	
8	85099B	JUMBO BAG RED WHITE SPOTTY	
9	85099B	JUMBO BAG RED WHITE SPOTTY	
10	47566	PARTY BUNTING	
11	85123A	WHITE HANGING HEART T-LIGHT HOLDER	
12	21755	LOVE BUILDING BLOCK WORD	

	recommendation_1_category	recommendation_2_product	\
1	HOME_DECOR	84879	
2	KITCHEN_FOOD_UTENSIL	47566	
3	HOME_DECOR	22469	
6	HOME_DECOR	84879	
7	FURNITURE_STORAGE	23298	
8	TEXTILES_CLOTHING	21754	
9	TEXTILES_CLOTHING	21754	
10	KITCHEN_FOOD_UTENSIL	22139	
11	HOME_DECOR	84879	
12	BEAUTY_PERSONAL	21790	

	recommendation_2_name	recommendation_2_category	\
1	ASSORTED COLOUR BIRD ORNAMENT	HOME_DECOR	
2	PARTY BUNTING	KITCHEN_FOOD_UTENSIL	

3	HEART OF WICKER SMALL	HOME_DECOR
6	ASSORTED COLOUR BIRD ORNAMENT	HOME_DECOR
7	SPOTTY BUNTING	GARDEN_OUTDOOR
8	HOME BUILDING BLOCK WORD	FURNITURE_STORAGE
9	HOME BUILDING BLOCK WORD	FURNITURE_STORAGE
10	RETRO SPOT TEA SET CERAMIC 11 PC	KITCHEN_FOOD_UTENSIL
11	ASSORTED COLOUR BIRD ORNAMENT	HOME_DECOR
12	VINTAGE SNAP CARDS	BEAUTY_PERSONAL

	recommendation_3_product	recommendation_3_name \
1	22469	HEART OF WICKER SMALL
2	21212	PACK OF 72 RETRO SPOT CAKE CASES
3	22138	BAKING SET 9 PIECE RETROSPOT
6	22469	HEART OF WICKER SMALL
7	21791	VINTAGE HEADS AND TAILS CARD GAME
8	21791	VINTAGE HEADS AND TAILS CARD GAME
9	21791	VINTAGE HEADS AND TAILS CARD GAME
10	22111	SCOTTIE DOG HOT WATER BOTTLE
11	22469	HEART OF WICKER SMALL
12	21915	RED HARMONICA IN BOX

	recommendation_3_category
1	HOME_DECOR
2	KITCHEN_FOOD_UTENSIL
3	HOME_DECOR
6	HOME_DECOR
7	TOYS_GAMES
8	TOYS_GAMES
9	TOYS_GAMES
10	KITCHEN_FOOD_UTENSIL
11	HOME_DECOR
12	BEAUTY_PERSONAL

Final Customer 360 Recommendation Table completed!

6. EXPORTING RECOMMENDATION TABLE:

- Exported to: datasets/customer_360_recommendations.csv

6.1 Summary and Recommendation